

U

Search for anything

CATALOGMY PROGRAMS

CONCEPTS

30. Create an AWS Redshift Cluster

31. Demo: Exercise 6: Building the S3 to Redshift DAG

32. Add Airflow Connections to AWS Redshift

33. Exercise 6: Build the S3 to Redshift DAG

34. Solution 6: Build the S3 to Redshift DAG

35. Conclusion

RESOURCES

LAUNCH CLOUD GATEWAY

31. Demo: Exercise 6: Building the S3 to Redshift DAG

32. Add Airflow Connections to AWS Redshift

33. Exercise 6: Build the S3 to Redshift DAG

34. Solution 6: Build the S3 to Redshift DAG

35. Conclusion

RESOURCES

LAUNCH CLOUD GATEWAY

31. Demo: Exercise 6: Building the S3 to Redshift DAG

32. Add Airflow Connections to AWS Redshift

33. Exercise 6: Build the S3 to Redshift DAG

34. Solution 6: Build the S3 to Redshift DAG

35. Conclusion

RESOURCES

LAUNCH CLOUD GATEWAY

31. Demo: Exercise 6: Building the S3 to Redshift DAG

32. Add Airflow Connections to AWS Redshift

33. Exercise 6: Build the S3 to Redshift DAG

34. Solution 6: Build the S3 to Redshift DAG

35. Conclusion

DATA ENGINEERING

AUTOMATE DATA PIPELINES

SEND PAGE FEEDBACK

?

10 minutes remaining

Data Pipelines

Add Airflow Connections to AWS Redshift

Here, we'll use Airflow's UI to connect to AWS Redshift

1. To go to the Airflow UI:

In a classroom workspace, run the script `/opt/airflow/start.sh`. Once you see the message "Airflow web server is ready" click on the blue **Access Airflow** button in the bottom right.

2. Click on the **Admin** tab and select **Connections**.

Airflow

DAGs

Data Profiling

Browse

Admin

Docs

About

DAGs

?

DAG

lesson1.exercise1

1 day, 0:00:00

Click on the Admin tab and select Connections.

3. Under **Connections**, select **Create**.

Airflow

DAGs

Data Profiling

Browse

Admin

Connections

List (36)

Create

With selected

Conn Id

airflow_db

Co

my

Under Connections, select Create.

On the create connection page, enter the following values:

• **Conn Id:** Enter `redshift`.

• **Conn Type:** Enter `Postgres`.

• **Host:** Enter the endpoint of your Redshift cluster, excluding the port and schema name at the end. You can find this by selecting your cluster in the **Clusters** page of the Amazon Redshift console. See where this is located in the screenshot below. IMPORTANT: Make sure to **NOT** include the port and schema name at the end of the Redshift endpoint string.

• **Schema:** Enter `dev`. This is the Redshift database you want to connect to.

• **Login:** Enter `awsuser`.

• **Password:** Enter the password you created when launching your Redshift cluster.

• **Port:** Enter `5439`. Once you've entered these values, select **Save**.

Amazon Redshift

Clusters

redshift-cluster-1

to Redshift DAG

✓

32. Add Airflow Connections to AWS Redshift

○

33. Exercise 6: Build the S3 to Redshift DAG

○

34. Solution 6: Build the S3 to Redshift DAG

○

35. Conclusion

RESOURCES ▲

LAUNCH CLOUD GATEWAY

✓

31. Demo: Exercise 6: Building the S3 to Redshift DAG

✓

32. Add Airflow Connections to AWS Redshift

○

33. Exercise 6: Build the S3 to Redshift DAG

✓

31. Demo: Exercise 6: Building the S3 to Redshift DAG

✓

32. Add Airflow Connections to AWS Redshift

○

33. Exercise 6: Build the S3 to Redshift DAG

○

34. Solution 6: Build the S3 to Redshift DAG

○

35. Conclusion

RESOURCES ▲

LAUNCH CLOUD GATEWAY

redshift-cluster-1 ▾				Actions ▾	Edit	Add partner integration	Query data ▾
General information 🔍							
Cluster identifier redshift-cluster-1	Status 🟢 Available	Node type dc2.large	Endpoint redshift-cluster-1.cnknr1caj9rz.us-west-2.redshift.amazonaws...				
Cluster namespace 54721f2c-7110-463f-9a67-79a7dc1c3b38	Date created February 14, 2022, 19:13 (UTC+05:30)	Number of nodes 1	JDBC URL jdbc:redshift://redshift-cluster-1.cnknr1caj9rz.us-west-2.reds...				
	Storage used 0.22% (0.35 of 160 GB used)	AQUA Not available	ODBC URL Driver=(Amazon Redshift (64-bit)); Server=redshift-cluster-1.cn...				

Redshift endpoint

Record was successfully created.

Connection [create]

List

Create

Conn Id

redshift

Conn Type

Postgres

Host

redshift-cluster-1.cro5lqt0mnm.us-west-2.redshift.amazonaws.com

Schema

dev

Login

awsuser

Password

Port

5439

Extra

Save

Save and Add Another

Save and Continue Editing

Cancel

Verify connection record was saved

Awesome! You're now all configured to run Airflow with Redshift.

WARNING: Remember to DELETE your cluster each time you are finished working to avoid large, unexpected costs.